

# Practice Problems 2

Jordan Lian

2/07/2021

1. (15 pts) The built-in dataset `USArrests` contains statistics about violent crime rates in the US States. Determine which states are outliers in terms of murders. Outliers, for the sake of this question, are defined as values that are more than 1.5 standard deviations from the mean.

```
# New Data Frame
df <- USArrests

# Reference Statistics
mean(df$Murder)

## [1] 7.788

1.5 * sd(df$Murder)

## [1] 6.533265

# Use row.names()
row.names(df)[abs(df$Murder - mean(df$Murder)) > (1.5*sd(df$Murder))]

## [1] "Florida"         "Georgia"        "Louisiana"      "Mississippi"
## [5] "North Dakota"    "South Carolina"
```

2. (15 pts) For the same dataset as in (1), is there a correlation between urban population and murder, i.e., as one goes up, does the other statistic as well? Comment on the strength of the correlation. Calculate the Pearson coefficient of correlation in R.

```
# Libraries
library(ggplot2)
library(ggpubr)

# Correlation Calculation
cor.test(df$Murder, df$UrbanPop)

##
## Pearson's product-moment correlation
##
## data: df$Murder and df$UrbanPop
## t = 0.48318, df = 48, p-value = 0.6312
## alternative hypothesis: true correlation is not equal to 0
```


```

## 95 percent confidence interval:
## -0.2128979 0.3413107
## sample estimates:
## cor
## 0.06957262

# Plot: Urban Population vs Murder
ggplot(df, aes(x=Murder, y=UrbanPop)) +
  geom_point() +
  geom_smooth() +
  stat_cor(method = "pearson", label.x = 7.5, label.y = 35) +
  ylab("Urban Population") +
  ggtitle("Urban Population vs Murder") +
  theme(plot.title = element_text(face = "bold", hjust = 0.5))

```

## ‘geom\_smooth()’ using method = ‘loess’ and formula ‘y ~ x’



The correlation value is 0.07, which shows basically no correlation between urban population and murder. From what I have read in my spare time, law enforcement and public policy have more to do with murder rather than urban population.

3. (3 x 10 pts) Based on the [data on the growth of mobile phone use in Brazil](#) (you will need to copy the data and create a CSV that you can load into R or use the `gsheet2tbl()` function from the `gsheet` package), forecast phone use for the next time period using a 2-year weighted moving average (with weights of 5 for the most recent year, and 2 for other), exponential smoothing (alpha of 0.4), and linear regression trendline.

I created separate data frames for each forecasting technique to show my results.

### Using `gsheet2tbl()` to load dataset into R, store link as variable URL

```
[1] "https://docs.google.com/spreadsheets/d/1tOnM9XceK4Ak8tzWQ2vDelWlJexzJiS3LbT6MN6_rW0/edit?usp=sharing"
```

```
# Import dataset Using gsheet2tbl()
library(gsheet)
brazil <- gsheet2tbl(url)
brazil

## # A tibble: 12 x 2
##       Year Subscribers
##   <dbl>     <dbl>
## 1     1    23188171
## 2     2    28745769
## 3     3    34880964
## 4     4    46373266
## 5     5    65605000
## 6     6    86210336
## 7     7    99918621
## 8     8   120980103
## 9     9   150641403
## 10   10   173959368
## 11   11   202944033
## 12   12        NA
```

### 2-year weighted moving average

$F_t = \frac{W_1*D_{t-1}+W_2*D_{t-2}}{W_1+W_2}$ , where  $F_t$  is the forecast,  $D_t$  is the demand (subscribers in this case), and  $W_t$  is the assigned weight to the demand value.

```
# New Data Frame, Initialize New Column
weight_MA <- brazil
weight_MA$Forecast <- rep(NA, 12)

# Store Weights Into a Vector
weights <- c(2, 5)

# Use weighted.mean() to get the weighted averages
for(i in 3:12){
  prev_2 <- c(brazil$Subscribers[(i-2):(i-1)])
  weight_MA$Forecast[i] <- weighted.mean(prev_2, weights)
}

# Forecast Data
weight_MA
```

```
## # A tibble: 12 x 3
##       Year Subscribers Forecast
##   <dbl>     <dbl>     <dbl>
```

```

## 1 23188171 NA
## 2 2 28745769 NA
## 3 3 34880964 27157884.
## 4 4 46373266 33128051.
## 5 5 65605000 43089751.
## 6 6 86210336 60110219.
## 7 7 99918621 80323097.
## 8 8 120980103 96001968.
## 9 9 150641403 114962537.
## 10 10 173959368 142166746.
## 11 11 202944033 167297092.
## 12 12 NA 194662700.

```

```

# Forecast for the Next Time Period
weight_MA$Forecast[12]

```

```

## [1] 194662700

```

### Exponential Smoothing (alpha = 0.4)

$F_{t+1} = F_t + \alpha(D_t - F_t)$ , where  $F_t$  is the forecast and  $D_t$  is the demand (subscribers in this case). Since this forecast uses previous forecast values to calculate the current forecast, we must initialize one value to start out.

```

# New Data Frame, Initialize New Column
expo_smooth <- brazil
expo_smooth$Forecast <- rep(NA, 12)

# Initialize First Forecast Value using Weighted Moving Average at Year 3
expo_smooth$Forecast[3] = weight_MA$Forecast[3]

# Get Forecast Values
for(i in 4:12){
  smoothing <- 0.4*(brazil$Subscribers[i-1] - expo_smooth$Forecast[i-1])
  expo_smooth$Forecast[i] <- expo_smooth$Forecast[i-1] + smoothing
}

# Forecast Data
expo_smooth

```

```

## # A tibble: 12 x 3
##       Year Subscribers   Forecast
##   <dbl>     <dbl>      <dbl>
## 1 1 23188171 NA
## 2 2 28745769 NA
## 3 3 34880964 27157884.
## 4 4 46373266 30247116.
## 5 5 65605000 36697576.
## 6 6 86210336 48260546.
## 7 7 99918621 63440462.
## 8 8 120980103 78031725.
## 9 9 150641403 95211076.

```

```
## 10     10    173959368 117383207.  
## 11     11    202944033 140013671.  
## 12     12          NA 165185816.
```

```
# Forecast for the Next Time Period  
expo_smooth$Forecast[12]
```

```
## [1] 165185816
```


## Linear Regression Trendline

```
# Get Regression Equation for Forecast Values  
lm(Subscribers ~ Year, data = brazil)
```

```
##  
## Call:  
## lm(formula = Subscribers ~ Year, data = brazil)  
##  
## Coefficients:  
## (Intercept)      Year  
## -15710760     18276748
```

```
# Plot the Data to Visualize  
ggscatter(brazil, x = "Year", y = "Subscribers", add = "reg.line") +  
  stat_cor(aes(label = paste(..rr.label.., ..p.label.., sep = "~`,"`~")))+  
  stat_regrline_equation(label.y = 1.8*100000000)
```

```
## `geom_smooth()` using formula 'y ~ x'
```



```

# New Data Frame, Initialize New Column
lin_reg <- brazil
lin_reg$Forecast <- rep(NA, 12)

# Get Forecast Values
for(i in 1:12){
  lin_reg$Forecast[i] <- -15710760 + (18276748 * i)
}

# Forecast Data
lin_reg

## # A tibble: 12 x 3
##   Year Subscribers Forecast
##   <dbl>     <dbl>    <dbl>
## 1     1     23188171  2565988
## 2     2     28745769  20842736
## 3     3     34880964  39119484
## 4     4     46373266  57396232
## 5     5     65605000  75672980
## 6     6     86210336  93949728
## 7     7     99918621 112226476
## 8     8    120980103 130503224
## 9     9    150641403 148779972
## 10    10   173959368 167056720
## 11    11   202944033 185333468

```

```
## 12      12          NA 203610216
```

```
# Forecast for the Next Time Period  
lin_reg$Forecast[12]
```

```
## [1] 203610216
```

4. (20 pts) Calculate the squared error for each model, i.e., use the model to calculate a forecast for each given time period and then the squared error. Finally, calculate the average (mean) squared error for each model. Which model has the smallest mean squared error (MSE)?

$E_t = F_t - D_t$ , where  $E_t$  is the Error value

$$MSE = \frac{1}{n} * \sum_{i=1}^n (E_i)^2$$

### Weighted Moving Average

```
# Create Error and Square Error Columns  
weight_MA>Error <- weight_MA$Forecast - weight_MA$Subscribers  
weight_MA$Sq_Error <- weight_MA>Error^2
```

```
# Square Error for Time Periods
```

```
weight_MA
```

```
## # A tibble: 12 x 5  
##   Year Subscribers Forecast     Error Sq_Error  
##   <dbl>     <dbl>    <dbl>     <dbl>    <dbl>  
## 1     1     23188171       NA       NA       NA  
## 2     2     28745769       NA       NA       NA  
## 3     3     34880964  27157884. -7723080.  5.96e13  
## 4     4     46373266  33128051. -13245215.  1.75e14  
## 5     5     65605000  43089751. -22515249.  5.07e14  
## 6     6     86210336  60110219. -26100117.  6.81e14  
## 7     7     99918621  80323097. -19595524.  3.84e14  
## 8     8    120980103  96001968. -24978135.  6.24e14  
## 9     9    150641403 114962537. -35678866.  1.27e15  
## 10    10   173959368 142166746. -31792622.  1.01e15  
## 11    11   202944033 167297092. -35646941.  1.27e15  
## 12    12          NA 194662700.       NA       NA
```

```
# MSE  
weight_MA_MSE <- mean(weight_MA$Sq_Error, na.rm = T)  
weight_MA_MSE
```

```
## [1] 6.650647e+14
```

### Exponential Smoothing

```

# Create Error and Square Error Columns
expo_smooth$Error <- expo_smooth$Forecast - expo_smooth$Subscribers
expo_smooth$Sq_Error <- expo_smooth$Error^2

# Square Error for Time Periods
expo_smooth

## # A tibble: 12 x 5
##   Year Subscribers Forecast      Error Sq_Error
##   <dbl>     <dbl>    <dbl>     <dbl>    <dbl>
## 1     1     23188171       NA       NA     NA
## 2     2     28745769       NA       NA     NA
## 3     3     34880964 27157884. -7723080.  5.96e13
## 4     4     46373266 30247116. -16126150.  2.60e14
## 5     5     65605000 36697576. -28907424.  8.36e14
## 6     6     86210336 48260546. -37949790.  1.44e15
## 7     7     99918621 63440462. -36478159.  1.33e15
## 8     8    120980103 78031725. -42948378.  1.84e15
## 9     9    150641403 95211076. -55430327.  3.07e15
## 10    10   173959368 117383207. -56576161.  3.20e15
## 11    11   202944033 140013671. -62930362.  3.96e15
## 12    12           NA 165185816.       NA     NA

# MSE
expo_smooth_MSE <- mean(expo_smooth$Sq_Error, na.rm = T)
expo_smooth_MSE

## [1] 1.778262e+15

```

## Linear Regression

```

# Create Error and Square Error Columns
lin_reg$Error <- lin_reg$Forecast - lin_reg$Subscribers
lin_reg$Sq_Error <- lin_reg$Error^2

# Square Error for Time Periods
lin_reg

## # A tibble: 12 x 5
##   Year Subscribers Forecast      Error Sq_Error
##   <dbl>     <dbl>    <dbl>     <dbl>    <dbl>
## 1     1     23188171 2565988 -20622183  4.25e14
## 2     2     28745769 20842736 -7903033  6.25e13
## 3     3     34880964 39119484  4238520  1.80e13
## 4     4     46373266 57396232 11022966  1.22e14
## 5     5     65605000 75672980 10067980  1.01e14
## 6     6     86210336 93949728  7739392  5.99e13
## 7     7     99918621 112226476 12307855  1.51e14
## 8     8    120980103 130503224  9523121  9.07e13
## 9     9    150641403 148779972 -1861431  3.46e12

```

```

## 10    10  173959368 167056720 -6902648  4.76e13
## 11    11  202944033 185333468 -17610565  3.10e14
## 12    12          NA 203610216          NA NA

# MSE
lin_reg_MSE <- mean(lin_reg$Sq_Error, na.rm = T)
lin_reg_MSE

## [1] 1.265347e+14

```

### Smallest Mean Squared Error (MSE)

Using the min() function, I found that the linear regression forecast produced the smallest MSE.

```
min(weight_MA_MSE, expo_smooth_MSE, lin_reg_MSE)
```

```
## [1] 1.265347e+14
```

5. (20 pts) Calculate a weighted average forecast by averaging out the three forecasts calculated in (3) with the following weights: 4 for trend line, 2 for exponential smoothing, 1 for weighted moving average. Remember to divide by the sum of the weights in a weighted average.

```

# New Data Frame, Initialize New Column
weight_AF <- brazil
weight_AF$Forecast <- rep(NA, 12)

# Store Weights Into a Vector
forecast_weights <- c(4, 2, 1)

# Use weighted.mean() to get the weighted averages
for(i in 3:12){
  forecast_vals <- c(lin_reg$Forecast[i], expo_smooth$Forecast[i], weight_MA$Forecast[i])
  weight_AF$Forecast[i] <- weighted.mean(forecast_vals, forecast_weights)
}

# Forecast Data
weight_AF

## # A tibble: 12 x 3
##       Year Subscribers   Forecast
##   <dbl>     <dbl>     <dbl>
## 1      1     23188171      NA
## 2      2     28745769      NA
## 3      3     34880964 33993084.
## 4      4     46373266 46172459.
## 5      5     65605000 59882403.
## 6      6     86210336 76061460.
## 7      7     99918621 93729989.
## 8      8    120980103 110582616.
## 9      9    150641403 128643511.
## 10     10    173959368 149308577.
## 11     11    202944033 169808330.
## 12     12          NA 191353599.

```