

Practice Problems 3

Jordan Lian

2/14/2021

1. Download the [data set for the tutorial](#).

```
library(tidyverse)

## -- Attaching packages ----

## v ggplot2 3.3.3      v purrr    0.3.4
## v tibble   3.0.3      v dplyr    1.0.2
## v tidyverse 1.1.2      v stringr  1.4.0
## v readr    1.3.1      vforcats 0.5.0

## -- Conflicts -----
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()   masks stats::lag()

origin_prostate <- read.csv("prostate_cancer.csv", stringsAsFactors = FALSE)
prostate <- origin_prostate
head(prostate)

##   id diagnosis_result radius texture perimeter area smoothness compactness
## 1  1             M     23     12     151  954    0.143    0.278
## 2  2             B      9     13     133 1326    0.143    0.079
## 3  3             M     21     27     130 1203    0.125    0.160
## 4  4             M     14     16      78  386    0.070    0.284
## 5  5             M      9     19     135 1297    0.141    0.133
## 6  6             B     25     25     83  477    0.128    0.170
##   symmetry fractal_dimension
## 1    0.242            0.079
## 2    0.181            0.057
## 3    0.207            0.060
## 4    0.260            0.097
## 5    0.181            0.059
## 6    0.209            0.076
```

2. Follow this [tutorial on applying kNN to prostate cancer detection](#) and implement all of the steps in an R Notebook. Make sure to explain each step and what it does. (*Note:* The data set provided as part of this assignment has been slightly modified from the one used in the tutorial, so small deviations in the result can be expected.)'

Step 1: Data Collection

See Question 1

Step 2: Preparing and Exploring the Data

```
# remove the first variable (id) from the dataset
prostate <- prostate[-1]

# get table of number of patients based on diagnosis
table(prostate$diagnosis_result)

## 
##   B   M
## 38 62

# Rename B and M to Benign and Malignant
prostate$diagnosis <- factor(prostate$diagnosis_result, levels = c("B", "M"),
                             labels = c("Benign", "Malignant"))
# Get percentage of B/M diagnoses
round(prop.table(table(prostate$diagnosis)) * 100, digits = 1)

## 
##      Benign Malignant
##          38       62
```

Normalize Numeric Data

```
normalize <- function(x) {
  return ((x - min(x)) / (max(x) - min(x)))
}
norm_prostate <- as.data.frame(lapply(prostate[2:9], normalize))
summary(norm_prostate$radius)

##      Min. 1st Qu. Median      Mean 3rd Qu.      Max.
## 0.0000  0.1875  0.5000  0.4906  0.7500  1.0000
```

Create test and training data set

```
# Split training and test dataset by 65/35
prostate_train <- norm_prostate[1:65,]
prostate_test <- norm_prostate[66:100,]

# Target variable is 'diagnosis_result', which is not in the training/test datasets
train_labels <- prostate[1:65, 1]
test_labels <- prostate[66:100, 1]
```

Step 3 - Train the model on data

```
library(class)
prc_test_pred <- knn(train = prostate_train,
                      test = prostate_test,
                      cl = train_labels,
                      k=10)
```

Step 4 - Evaluate the model performance

```
library(gmodels)
CrossTable(x = test_labels, y = prc_test_pred, prop.chisq = FALSE)

##
##
##      Cell Contents
## |-----|
## |           N |
## |           N / Row Total |
## |           N / Col Total |
## |           N / Table Total |
## |-----|
## 
## 
## Total Observations in Table:  35
##
##
##          | prc_test_pred
## test_labels |      B |      M | Row Total |
## -----|-----|-----|-----|
##       B |      6 |     13 |      19 |
##       | 0.316 | 0.684 | 0.543 |
##       | 0.857 | 0.464 |      |
##       | 0.171 | 0.371 |      |
## -----|-----|-----|-----|
##       M |      1 |     15 |      16 |
##       | 0.062 | 0.938 | 0.457 |
##       | 0.143 | 0.536 |      |
##       | 0.029 | 0.429 |      |
## -----|-----|-----|-----|
## Column Total |      7 |     28 |      35 |
##       | 0.200 | 0.800 |      |
## -----|-----|-----|-----|
## 
```

Step 5 - Improve the performance of the model, change the k-value

```

# k=8
prc_test_pred <- knn(train = prostate_train,
                      test = prostate_test,
                      cl = train_labels,
                      k = 8)
CrossTable(x = test_labels, y = prc_test_pred, prop.chisq = FALSE)

```

```

##
##
##      Cell Contents
## |-----|
## |           N |
## |   N / Row Total |
## |   N / Col Total |
## |   N / Table Total |
## |-----|
##
##
## Total Observations in Table:  35
##
##
##          | prc_test_pred
## test_labels |      B |      M | Row Total |
## -----|-----|-----|-----|
##       B |     9 |    10 |     19 |
##       | 0.474 | 0.526 | 0.543 |
##       | 1.000 | 0.385 |      |
##       | 0.257 | 0.286 |      |
## -----|-----|-----|-----|
##       M |     0 |    16 |     16 |
##       | 0.000 | 1.000 | 0.457 |
##       | 0.000 | 0.615 |      |
##       | 0.000 | 0.457 |      |
## -----|-----|-----|-----|
## Column Total |     9 |    26 |     35 |
##       | 0.257 | 0.743 |      |
## -----|-----|-----|-----|
##
##

```

```

# k=9
prc_test_pred <- knn(train = prostate_train,
                      test = prostate_test,
                      cl = train_labels,
                      k = 9)
CrossTable(x = test_labels, y = prc_test_pred, prop.chisq = FALSE)

```

```

##
##
##      Cell Contents
## |-----|
## |           N |

```

```

## | N / Row Total |
## | N / Col Total |
## | N / Table Total |
## |-----|
## 
## 
## Total Observations in Table: 35
## 
## 
## | prc_test_pred
## test_labels | B | M | Row Total |
## -----|-----|-----|-----|
## B | 8 | 11 | 19 |
## | 0.421 | 0.579 | 0.543 |
## | 1.000 | 0.407 | |
## | 0.229 | 0.314 | |
## -----|-----|-----|
## M | 0 | 16 | 16 |
## | 0.000 | 1.000 | 0.457 |
## | 0.000 | 0.593 | |
## | 0.000 | 0.457 | |
## -----|-----|-----|
## Column Total | 8 | 27 | 35 |
## | 0.229 | 0.771 | |
## -----|-----|-----|
## 
## 

# k=11
prc_test_pred <- knn(train = prostate_train,
                      test = prostate_test,
                      cl = train_labels,
                      k = 11)
CrossTable(x = test_labels, y = prc_test_pred, prop.chisq = FALSE)

## 
## 
## Cell Contents
## |-----|
## | N |
## | N / Row Total |
## | N / Col Total |
## | N / Table Total |
## |-----|
## 
## 
## Total Observations in Table: 35
## 
## 
## | prc_test_pred
## test_labels | B | M | Row Total |
## -----|-----|-----|-----|
## B | 5 | 14 | 19 |
## | 0.263 | 0.737 | 0.543 |

```

```

##          |    1.000 |    0.467 |
##          |    0.143 |    0.400 |
## -----
##      M |      0 |      16 |      16 |
##          |    0.000 |    1.000 |    0.457 |
##          |    0.000 |    0.533 |
##          |    0.000 |    0.457 |
## -----
## Column Total |      5 |     30 |     35 |
##          |    0.143 |    0.857 |
## -----
##
```

```

# k=12
prc_test_pred <- knn(train = prostate_train,
                      test = prostate_test,
                      cl = train_labels,
                      k = 12)
CrossTable(x = test_labels, y = prc_test_pred, prop.chisq = FALSE)

```

```

##
##
##      Cell Contents
## |-----|
## |           N |
## |           N / Row Total |
## |           N / Col Total |
## |           N / Table Total |
## |-----|
##
##
## Total Observations in Table: 35
##
##
##          | prc_test_pred
##  test_labels |      B |      M | Row Total |
## -----|-----|-----|-----|
##      B |      5 |     14 |      19 |
##          | 0.263 | 0.737 | 0.543 |
##          | 1.000 | 0.467 |
##          | 0.143 | 0.400 |
## -----|-----|-----|-----|
##      M |      0 |      16 |      16 |
##          | 0.000 | 1.000 | 0.457 |
##          | 0.000 | 0.533 |
##          | 0.000 | 0.457 |
## -----|-----|-----|-----|
## Column Total |      5 |     30 |     35 |
##          |    0.143 |    0.857 |
## -----|-----|-----|-----|
##
```

k=9 produced the least amount of false positives and most accurate results overall, so I would use k=9 for kNN for this dataset.

3. Once you have complete the tutorial, try another kNN implementation from another package, such as the **caret** package. Compare the accuracy of the two implementations.

```
library(caret)

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##     lift
```

Sampling

```
# Split data as training and test set. Using createDataPartition() function from caret
indxTrain <- createDataPartition(y = prostate$diagnosis_result, p = 0.75, list = FALSE)
training <- prostate[indxTrain,]
testing <- prostate[-indxTrain,]

# Check distribution in original data and partitioned data
prop.table(table(prostate$diagnosis_result)) * 100
```

```
##
##     B     M
## 38   62

prop.table(table(testing$diagnosis_result)) * 100
```

```
##
##     B     M
## 37.5 62.5

prop.table(table(prostate$diagnosis_result)) * 100
```

```
##
##     B     M
## 38   62
```

Preprocessing

```

trainX <- training[,names(training) != "diagnosis_result"]
preProcValues <- preProcess(x = trainX, method = c("center", "scale"))
preProcValues

```

```

## Created from 76 samples and 9 variables
##
## Pre-processing:
##   - centered (8)
##   - ignored (1)
##   - scaled (8)

```

Training and train control

```

set.seed(400)
ctrl <- trainControl(method="repeatedcv", repeats = 3)
knnFit <- train(diagnosis_result ~ .,
                  data = training,
                  method = "knn",
                  trControl = ctrl,
                  preProcess = c("center","scale"),
                  tuneLength = 20)

# kNNFit output
knnFit

```

```

## k-Nearest Neighbors
##
## 76 samples
## 9 predictor
## 2 classes: 'B', 'M'
##
## Pre-processing: centered (9), scaled (9)
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 69, 69, 68, 68, 68, 69, ...
## Resampling results across tuning parameters:
##
##     k    Accuracy   Kappa
##     5    0.9494048  0.8785822
##     7    0.9577381  0.8976298
##     9    0.9619048  0.9071536
##    11    0.9577381  0.8976298
##    13    0.9535714  0.8881060
##    15    0.9440476  0.8678162
##    17    0.9398810  0.8568271
##    19    0.9357143  0.8473033
##    21    0.9488095  0.8779611
##    23    0.9404762  0.8589135
##    25    0.9279762  0.8274116
##    27    0.9184524  0.8071218
##    29    0.9095238  0.7874530

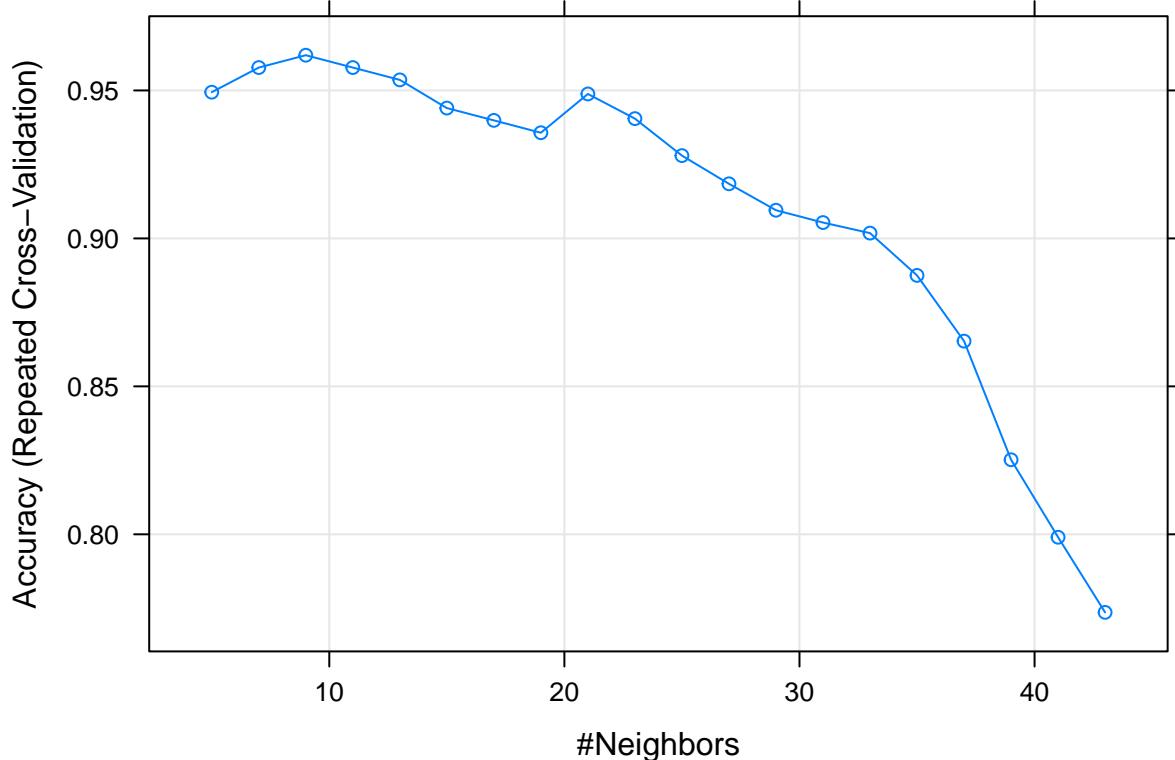
```

```

##   31  0.9053571  0.7764640
##   33  0.9017857  0.7627994
##   35  0.8875000  0.7314424
##   37  0.8652778  0.6790614
##   39  0.8251984  0.5763410
##   41  0.7990079  0.5033357
##   43  0.7736111  0.4469255
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 9.

```

```
#Plot yields Number of Neighbors Vs accuracy (based on repeated cross validation)
plot(knnFit)
```



4. Try the `confusionMatrix` function from the **caret** package to determine the accuracy of both algorithms.

```
knnPredict <- predict(knnFit, newdata = testing)

# Get the confusion matrix to see accuracy value and other parameter values
confusionMatrix(knnPredict, as.factor(testing$diagnosis_result))
```

```
## Confusion Matrix and Statistics
##
##          Reference
```

```

## Prediction B M
##      B 7 0
##      M 2 15
##
##          Accuracy : 0.9167
##      95% CI : (0.73, 0.9897)
##      No Information Rate : 0.625
##      P-Value [Acc > NIR] : 0.001448
##
##          Kappa : 0.814
##
##  Mcnemar's Test P-Value : 0.479500
##
##          Sensitivity : 0.7778
##          Specificity : 1.0000
##          Pos Pred Value : 1.0000
##          Neg Pred Value : 0.8824
##          Prevalence : 0.3750
##          Detection Rate : 0.2917
##          Detection Prevalence : 0.2917
##          Balanced Accuracy : 0.8889
##
##          'Positive' Class : B
##

mean(knnPredict == testing$diagnosis_result)

## [1] 0.9166667

# Verify 2 class summary function
ctrl <- trainControl(method="repeatedcv",
                      repeats = 3,
                      classProbs = TRUE,
                      summaryFunction = twoClassSummary)
knnFit <- train(diagnosis_result ~ .,
                 data = training,
                 method = "knn",
                 trControl = ctrl,
                 preProcess = c("center","scale"),
                 tuneLength = 20)

knnFit

## k-Nearest Neighbors
##
## 76 samples
## 9 predictor
## 2 classes: 'B', 'M'
##
## Pre-processing: centered (9), scaled (9)
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 68, 68, 68, 69, 70, 69, ...
## Resampling results across tuning parameters:
```

```

##  

##   k    ROC      Sens      Spec  

##   5  0.9683333  0.9055556  1  

##   7  0.9861111  0.8944444  1  

##   9  0.9933333  0.8944444  1  

##  11  0.9888889  0.8944444  1  

##  13  0.9905556  0.8722222  1  

##  15  0.9816667  0.8722222  1  

##  17  0.9805556  0.8388889  1  

##  19  0.9777778  0.8611111  1  

##  21  0.9722222  0.8611111  1  

##  23  0.9722222  0.8500000  1  

##  25  0.9711111  0.8388889  1  

##  27  0.9711111  0.8166667  1  

##  29  0.9744444  0.7833333  1  

##  31  0.9766667  0.7833333  1  

##  33  0.9700000  0.7888889  1  

##  35  0.9683333  0.8055556  1  

##  37  0.9711111  0.7111111  1  

##  39  0.9683333  0.6000000  1  

##  41  0.9655556  0.4666667  1  

##  43  0.9666667  0.3888889  1  

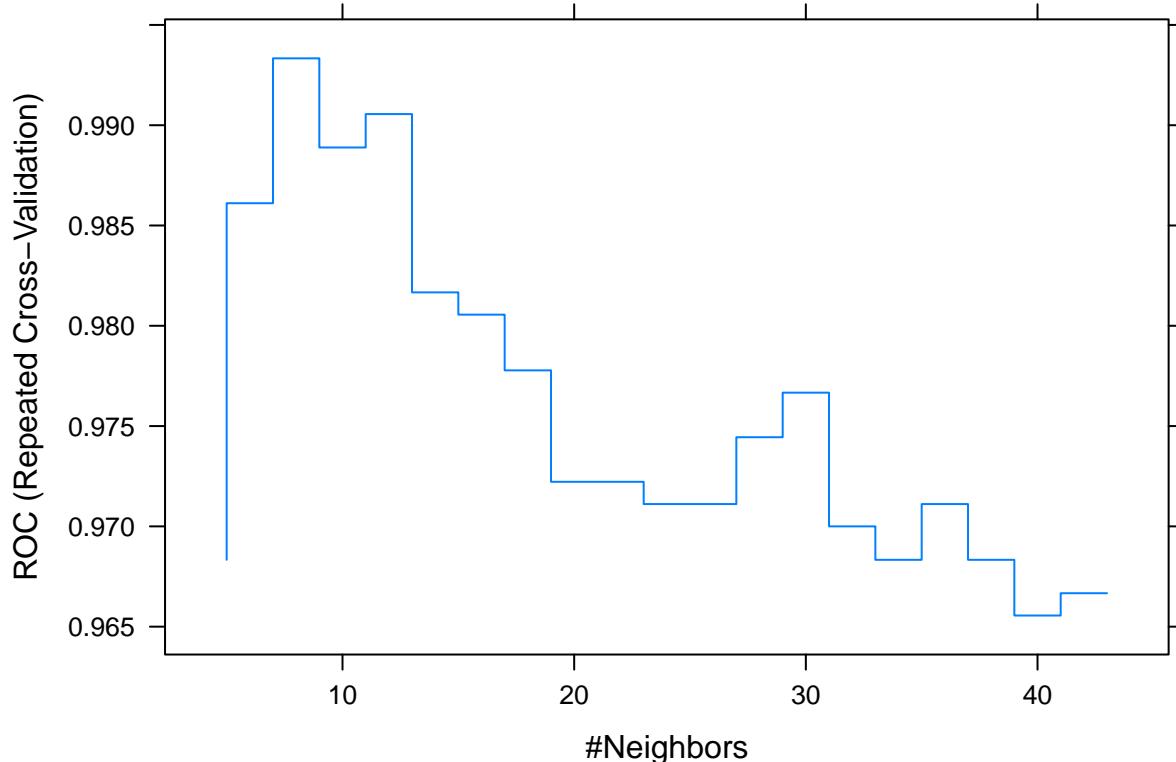
##  

## ROC was used to select the optimal model using the largest value.  

## The final value used for the model was k = 9.

```

```
plot(knnFit, print.thres = 0.5, type="S")
```



```

knnPredict <- predict(knnFit,newdata = testing)

# Get the confusion matrix to see accuracy value and other parameter values
confusionMatrix(knnPredict, as.factor(testing$diagnosis_result))

```

```

## Confusion Matrix and Statistics
##
##             Reference
## Prediction   B   M
##           B 7  0
##           M 2 15
##
##                  Accuracy : 0.9167
##                  95% CI : (0.73, 0.9897)
##      No Information Rate : 0.625
##      P-Value [Acc > NIR] : 0.001448
##
##                  Kappa : 0.814
##
## McNemar's Test P-Value : 0.479500
##
##                  Sensitivity : 0.7778
##                  Specificity  : 1.0000
##      Pos Pred Value : 1.0000
##      Neg Pred Value : 0.8824

```

```
##           Prevalence : 0.3750
##           Detection Rate : 0.2917
##   Detection Prevalence : 0.2917
##           Balanced Accuracy : 0.8889
##
##           'Positive' Class : B
##
mean(knnPredict == testing$diagnosis_result)

## [1] 0.9166667
```